

# Pydantic RAG

---

Agentic RAG chatbot built with Pydantic AI, Weaviate vector database, and local Ollama inference. Features hybrid search, conversation memory, and multiple RAG modes.

## Features

- **Hybrid Search** - Combines BM25 keyword search with semantic vector search (configurable alpha)
- **3 RAG Modes** - Auto (agent decides), Force (always search), Disabled (plain chat)
- **Conversation Memory** - Multi-turn conversations with full context via `message_history`
- **Token Tracking** - Real-time token usage display with context limit warnings
- **Local Inference** - GPU-accelerated LLM and embedding generation via Ollama

## Tech Stack

Component	Technology
Agent Framework	<a href="#">Pydantic AI</a>
Vector Database	<a href="#">Weaviate</a> (with <code>text2vec-ollama</code> )
LLM Inference	<a href="#">Ollama</a> ( <code>llama3.2</code> )
Embeddings	<code>nomic-embed-text</code> (768 dimensions)
Web UI	<a href="#">Gradio</a>
Orchestration	Docker Compose

## Prerequisites

PROF

- Docker with Compose v2
- NVIDIA GPU with CUDA drivers (for GPU inference)
- NVIDIA Container Toolkit (`nvidia-docker`)

## Quick Start

### 1. Clone and start services:

```
git clone <repo-url>
cd pydantic-rag
docker compose up -d
```

### 2. Wait for model downloads (first run only):

```
# Watch Ollama logs until models are ready
docker logs -f ollama
```

### 3. Access the UI at <http://localhost:7860>

## Document Ingestion

Place documents in [data/documents/](#) and run the ingestion script:

```
# Create documents directory
mkdir -p data/documents

# Add your documents (supports .txt, .md, .pdf, .py, .js, .ts, .c, .cpp,
.cu, .h)
cp /path/to/your/docs/* data/documents/

# Install dependencies and run ingestion
pip install weaviate-client pypdf
python scripts/ingest.py --name "my project"
```

### Ingestion options:

```
python scripts/ingest.py --help
python scripts/ingest.py --reset                                #
Reset collection only (no ingestion)
python scripts/ingest.py --name "my project" --reset           #
Delete and recreate collection, then ingest
python scripts/ingest.py --name "eu ai regulations" --extensions .pdf  #
Only PDFs with label
python scripts/ingest.py --name "docs" --documents-dir ./my-docs   #
Custom source folder
python scripts/ingest.py --name "code" --extensions .py,.md        # Only
Python and Markdown files
python scripts/ingest.py --name "config" --extensions .yml,Dockerfile #
Extensions and exact filenames
```

PROF

**Note:** The `--name` option is required when ingesting documents. Use `--reset` alone to recreate the schema without ingesting.

Documents are chunked (800 tokens, 200 overlap) and embedded automatically by Weaviate's `text2vec-ollama` module.

## Usage

### RAG Modes

Mode	Behavior
<b>Auto</b>	Agent decides when to search documents based on the question
<b>Force</b>	Always searches documents before answering
<b>Disabled</b>	Plain chat without document retrieval

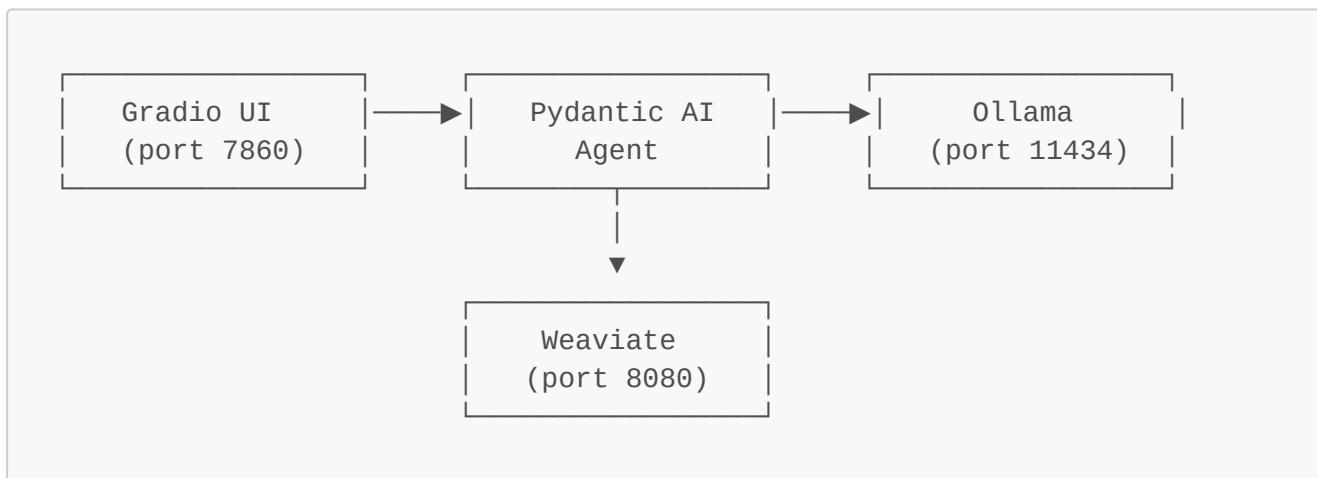
## Chat Interface

1. Select a RAG mode
2. Type your question and press Enter
3. View token usage in the top-right display
4. Click "Reset Chat" to clear conversation history

## Status Checks

Use the "Check Ollama" and "Check Weaviate" buttons to verify connections and see available models/collections.

## Architecture



PROF

## Configuration

Environment variables (set in `docker-compose.yml`):

Variable	Default	Description
OLLAMA_BASE_URL	<code>http://ollama:11434</code>	Ollama API endpoint
WEAVIATE_URL	<code>http://weaviate:8080</code>	Weaviate endpoint
CHAT_MODEL	<code>llama3.2</code>	LLM for chat
EMBED_MODEL	<code>nomic-embed-text</code>	Model for embeddings

## Troubleshooting

**Models not loading:** Check Ollama logs with `docker logs ollama`. First startup downloads ~2GB of models.

**Weaviate connection errors:** Ensure Weaviate is healthy with `docker compose ps`. The app will show connection status.

**Out of GPU memory:** llama3.2 (3B) requires ~4GB VRAM. For larger models, adjust `CHAT_MODEL` or use CPU inference.